

Trent Limited FAQ Chatbot – Report

Roll No.: L05	Name: k John Austin
Program : B.tech CSEDs	Sem/Year : VII / 4 th Year

Introduction

Trent Limited operates brands such as Westside, Zudio, Samoh, and Star. Stakeholders frequently seek concise, accurate answers about the company and its brands. This project delivers a Streamlit chatbot that uses only FAQs_questions.pdf as its knowledge base, ensuring responses remain within the document's scope.

Objectives

- Build a chatbot that replicates the existing app's flow and API usage but restricts context to FAQs_questions.pdf only.
- Provide transparent citations of the PDF snippets used for each answer.
- Maintain a simple, dependable retrieval pipeline for fast responses and easy deployment.

Dataset/Source Description

- Source: FAQs_questions.pdf (bundled with the app)
- Structure: Free-form text extracted per page using PyPDF2.PdfReader.
- Preprocessing: Minimal text cleaning; pages are treated as atomic context units.

Methodology

Retrieval-augmented generation (RAG) with:

- Vectorization: CountVectorizer(max_features=5000, stop_words='english')
- Weighting: TfidfTransformer(norm='l2', use_idf=True)
- Similarity: cosine_similarity between query vector and page TF-IDF vectors
- Generation: OpenRouter Chat Completions (openai/gpt-3.5-turbo by default)

Rationale: TF-IDF is light-weight, explainable, and effective for short FAQs; it avoids external indexes and reduces operational complexity for coursework and small deployments.

System Design & Architecture

- Ingestion: PyPDF2 extracts text per page from FAQs_questions.pdf.
- Indexing: Pages are vectorized and transformed into a TF-IDF matrix (cached via st.cache_resource).
- Retrieval: Top-k pages selected by cosine similarity to the user query.
- Prompting: Retrieved snippets are concatenated into the user prompt; the system prompt instructs grounding in the PDF.

- Generation: The LLM produces an answer with the provided context.
- Transparency: The UI displays the snippets (page numbers and similarity scores) used.

Component summary:

- `load_faq_artifacts(pdf_path)`: PDF ingestion and TF-IDF matrix build
- `retrieve_context(query, vectorizer, tfidf_matrix, pages, k)`: top-k retrieval
- `build_prompt(question, snippets)`: system and user prompts
- `call_openrouter(api_key, model, system_prompt, user_prompt)`: API call

Implementation Details

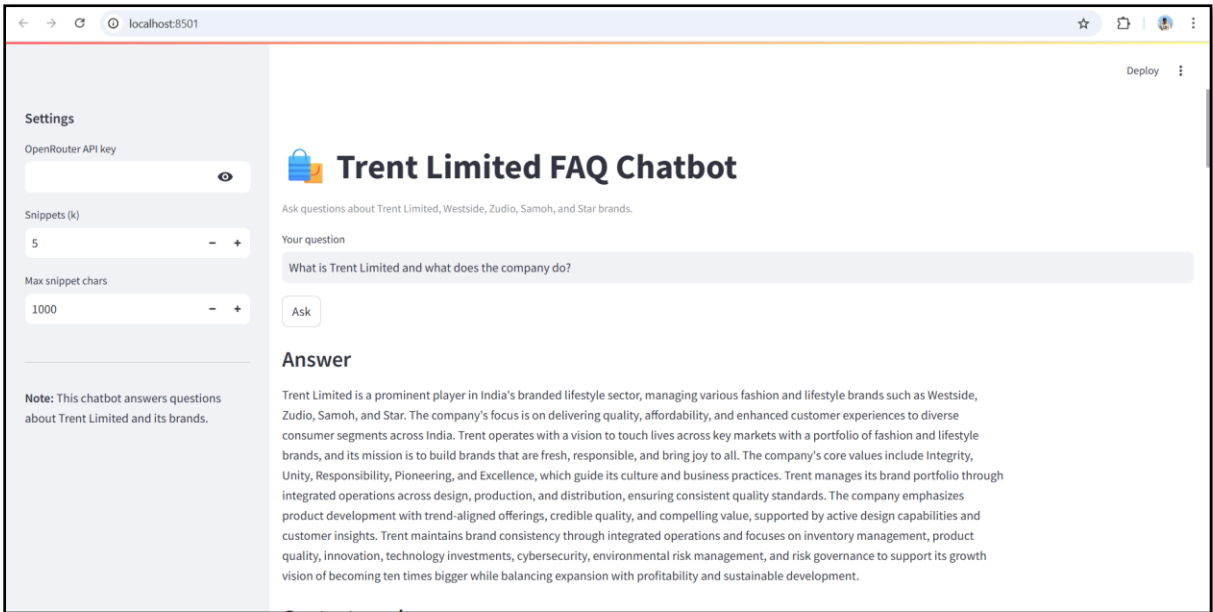
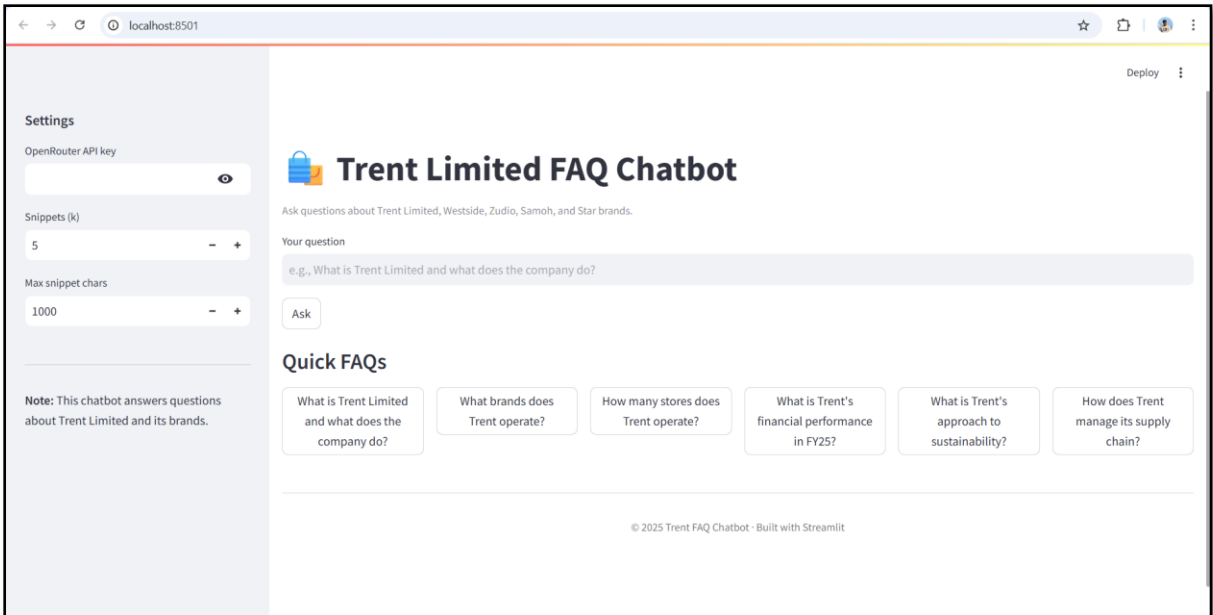
- Language & Frameworks: Python, Streamlit
- Dependencies: streamlit, scikit-learn, numpy, requests, PyPDF2
- Model API: OpenRouter (<https://openrouter.ai/api/v1/chat/completions>)
- Model: openai/gpt-3.5-turbo (configurable)
- Key Management: sidebar input, environment variable `OPENROUTER_API_KEY`, or Streamlit secrets
- Performance: k and snippet length tunable in sidebar; TF-IDF artifacts cached

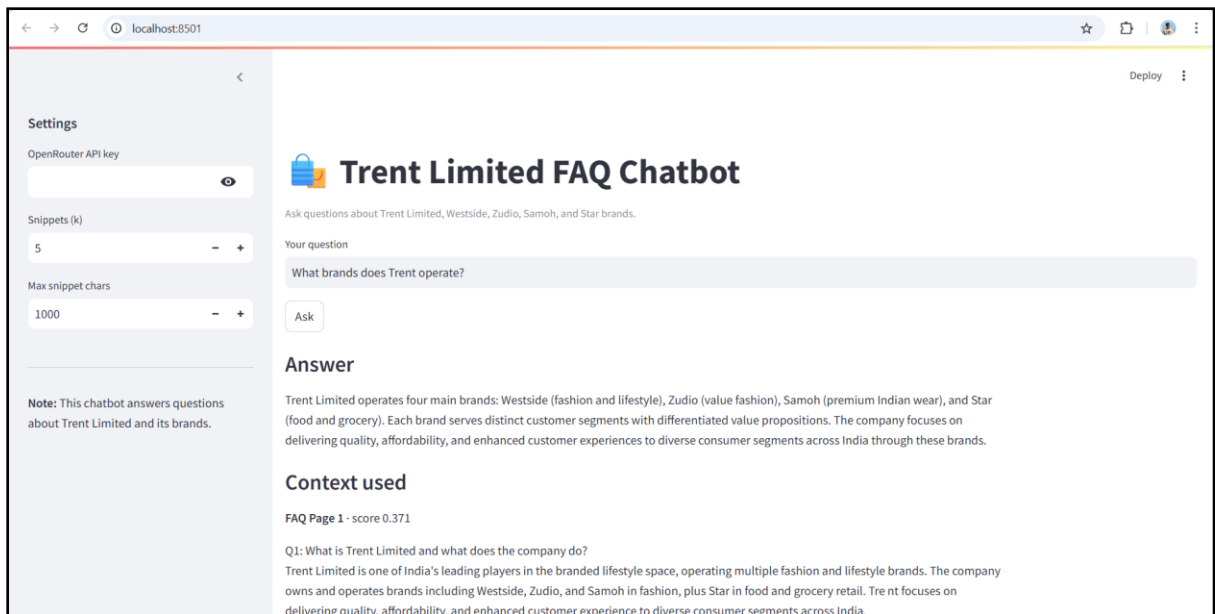
Key constraints enforced in code:

- Only `FAQs_questions.pdf` is read for knowledge.
- If the answer is not in the PDF, the bot states this explicitly and suggests relevant sections.

Results & Observations

- The TF-IDF retriever reliably surfaced on-topic pages for most FAQ-style questions.
- Concise, specific queries improved retrieval precision compared to very broad prompts.
- The transparency panel (snippets + scores) helped verify that answers were grounded in the PDF.





Setup & Configuration

```
python -m venv .venv
```

```
.venv\Scripts\activate
```

```
pip install -r requirements.txt
```

API key options:

- Sidebar input (masked)
- Environment variable: OPENROUTER_API_KEY
- Streamlit secrets: create .streamlit/secrets.toml with:

```
OPENROUTER_API_KEY = "sk-or-v1-..."
```

How to Run

```
streamlit run trent_faq_pdf_chatbot.py
```

Open the provided local URL. If needed, paste your API key in the sidebar. Ask a question and review the “Context used” panel for the supporting PDF snippets.

Conclusion

This project delivers a focused, PDF-grounded chatbot for Trent/Zudio that mirrors the original app’s API and UX while strictly constraining knowledge to FAQs_questions.pdf. The lightweight TF-IDF retriever provides dependable performance for FAQ-style questions and clear transparency of sources, satisfying the assignment’s requirements.

References

- Streamlit documentation — <https://docs.streamlit.io>
- scikit-learn: CountVectorizer, TfidfTransformer, cosine similarity — <https://scikit-learn.org>
- OpenRouter API — <https://openrouter.ai>

- PyPDF2 — <https://pypi.org/project/PyPDF2>

Conclusion

This project delivers a focused, PDF-grounded chatbot for Trent/Zudio that mirrors the original app's API and UX while strictly constraining knowledge to `FAQs_questions.pdf`. The lightweight TF-IDF retriever provides dependable performance for FAQ-style questions and clear transparency of sources, satisfying the assignment's requirements.